

ADVANCED IO

Tong Van Van, SoICT, HUST

Content

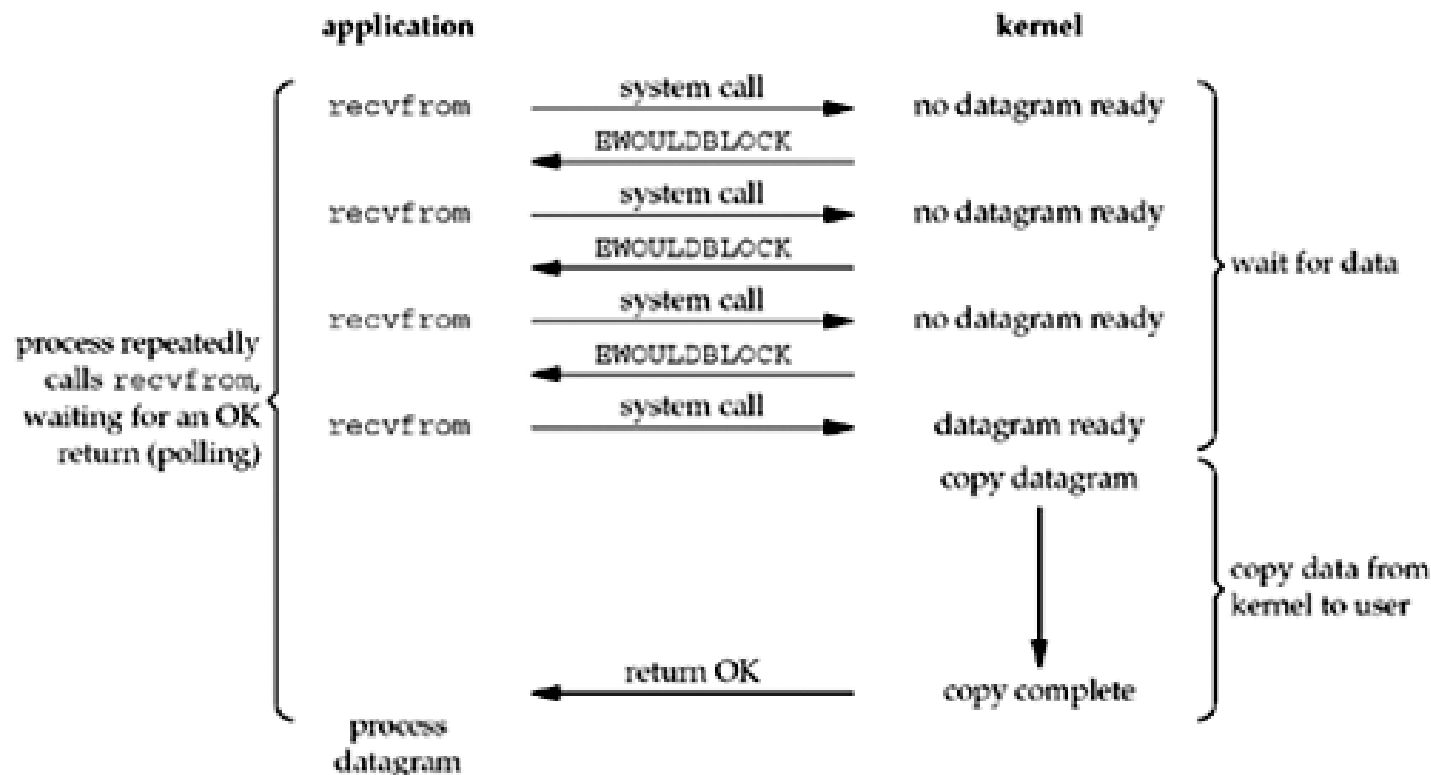
- Non-blocking IO
- Signal-driven I/O
- Some advanced I/O functions

Blocking I/O

- By default, sockets are blocking: when a socket call cannot be completed immediately, the process is put to sleep, waiting for the condition to be true
- Input functions: `recv()`, `recvfrom()`, etc
 - Blocks until some data arrives
- Output function: `send()`, `sendto()`, etc
 - TCP: blocks until there is free space in sending buffer
 - UDP: block on some systems due to the buffering and flow control
- Accepting incoming connections: `accept()`
 - Blocks until a new connection is available
- Initiating outgoing connections: `connect()`
 - Blocks until the client receives the ACK of its SYN

Non-blocking I/O Model

- Non-blocking I/O model: I/O function returns immediately
- If there is no data to return, so the kernel immediately returns an error of EWOULDBLOCK instead



Non-blocking I/O: use `fcntl()`

```
#include <fcntl.h>
int fcntl(int fd, int cmd, ... /* int arg */ );
```

- Perform the file control operations described below on open files
- Parameter:
 - [IN] `fd`: the file descriptor
 - [IN] `cmd`: the control operation
 - The 3rd argument according to `cmd`
- Return:
 - Return -1 on error
 - Otherwise, return others depending on `cmd`

Non-blocking I/O: use `fcntl()`

- Set non-blocking mode

```
int flags;
/* Get the file status flags and file access modes */
if ((flags = fcntl(fd, F_GETFL, 0)) < 0)
    perror("F_GETFL error");
/* Set a socket as nonblocking */
if (fcntl(fd, F_SETFL, flags | O_NONBLOCK) < 0)
    perror("F_SETFL error");
```

- Turn off non-blocking mode

```
int flags;
/* Get the file status flags and file access modes */
if ((flags = fcntl(fd, F_GETFL, 0)) < 0)
    perror("F_GETFL error");
/* Turn off non-blocking mode on socket */
if (fcntl(fd, F_SETFL, flags & ~O_NONBLOCK) < 0)
    perror("F_SETFL error");
```

Non-blocking I/O: use ioctl()

```
#include <sys/ioctl.h>
int ioctl(int fd, int request, ... /* void arg */ );
```

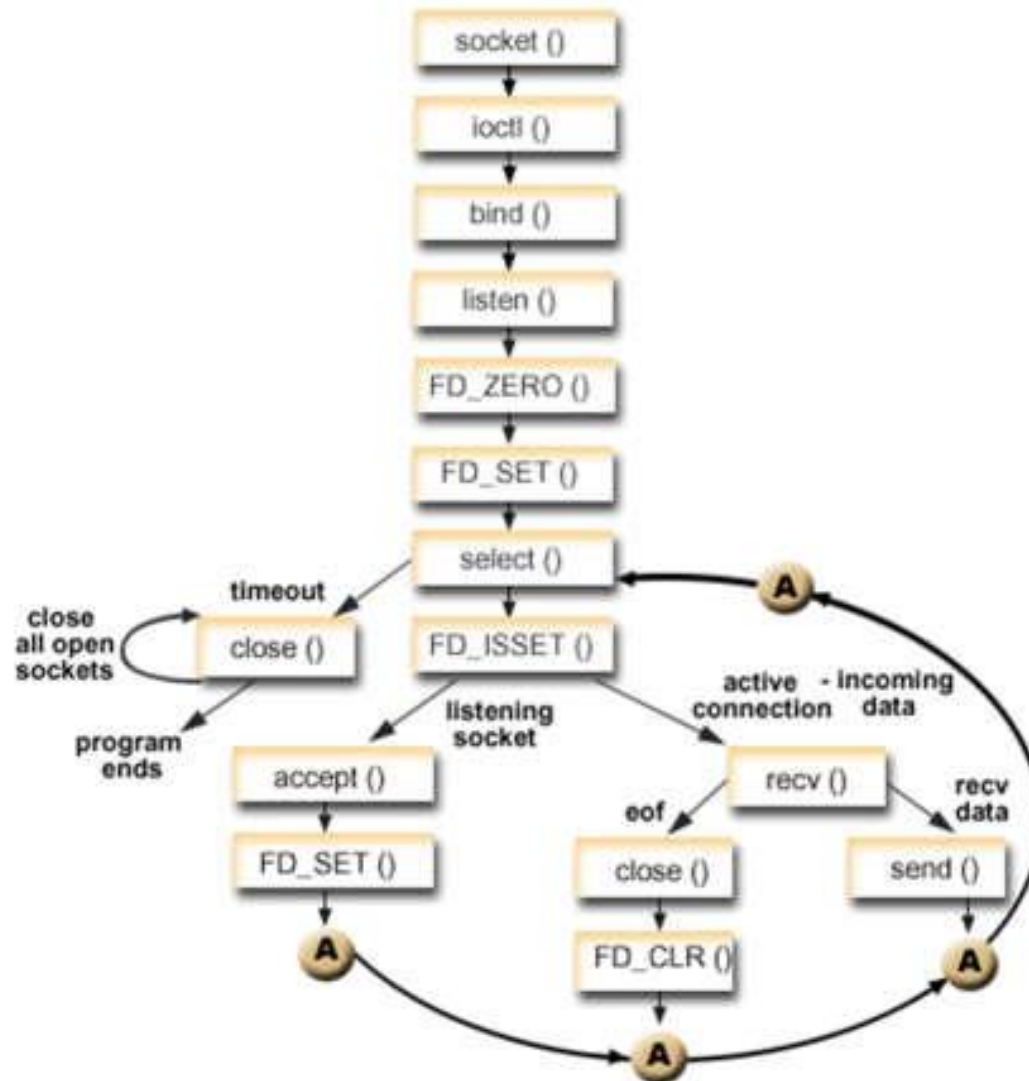
- Manipulates the underlying device parameters of special files and control operating characteristics of files
- Parameters
 - [IN]fd: the file descriptor
 - [IN]request: device-dependent request code
 - The 3rd argument according to request
- Return:
 - 0 if succeed
 - -1 if error

```
int on = 1;
/* Set a socket as nonblocking */
ioctl(fd, FIONBIO, (char *)&on);
on = 0;
/* Turn off non-blocking mode on socket */
ioctl(fd, FIONBIO, (char *)&on);
```

Non-blocking I/O: process return value

```
//call I/O functions
if (ret < 0) {
    //Error on I/O operation
    if (errno != EWOULDBLOCK)
    {
    }
    else if (ret == 0) {
        //Connection terminated normally
    }
    else if {
        //I/O operation is successful
    }
}
```


Non-blocking I/O: Example

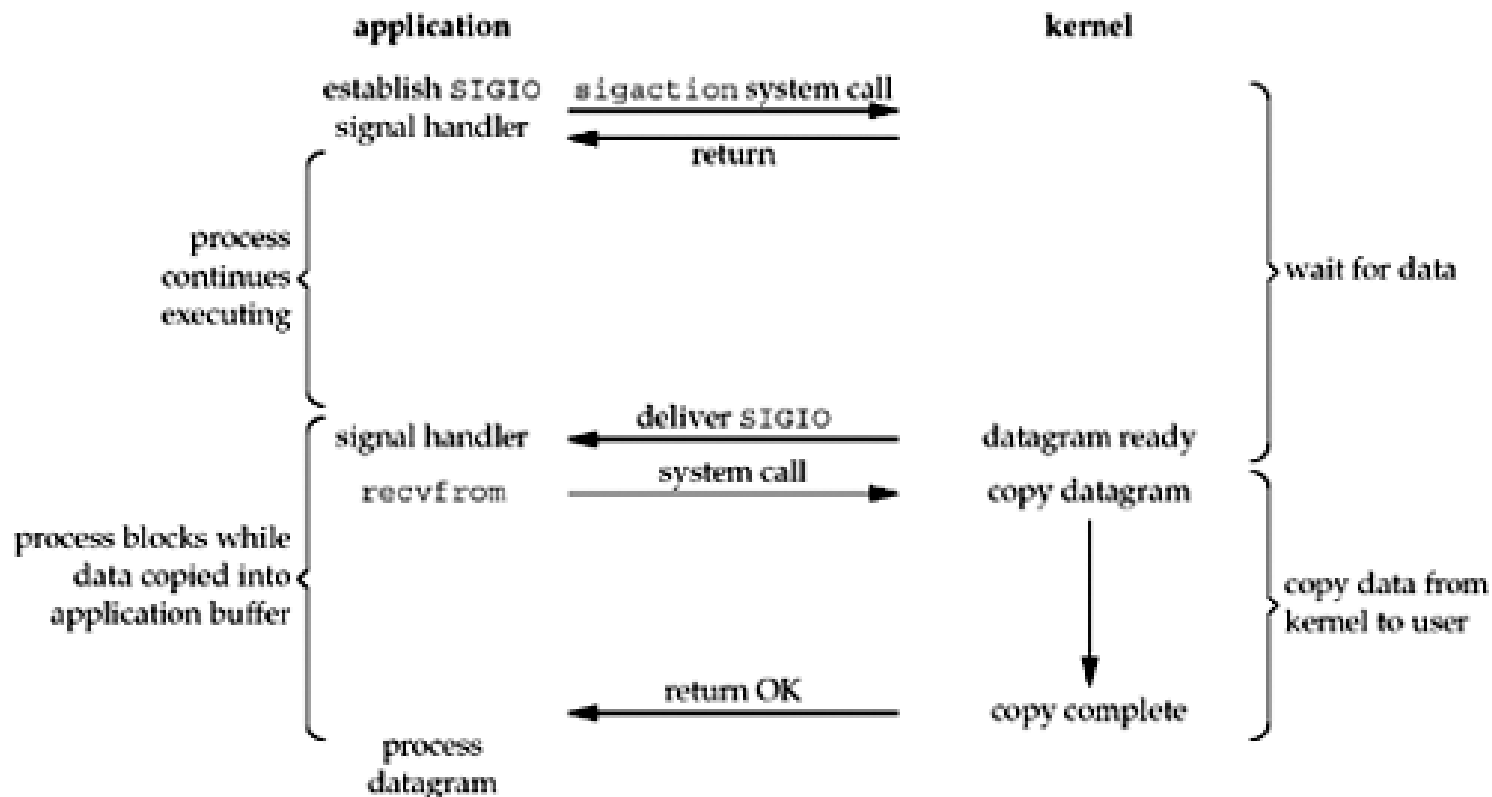


SIGNAL-DRIVEN I/O

Tong Van Van, SoICT, HUST

Signal-driven I/O

- Use signals, telling the kernel to notify app with the SIGIO signal when the descriptor is ready



Signal-driven I/O: 3 steps

1. A signal handler must be established for the SIGIO signal.
2. Assign a process to receive the SIGIO signal
`fcntl(fd, F_SETOWN, process_id)`
3. Enable signal-driven I/O on socket
 - Turn on asynchronous mode
 - Turn on non-blocking mode
- The importance is determining what conditions cause SIGIO to be generated for the socket owner

Signal-driven I/O: use `fcntl()`

- Enable signal-driven I/O on socket

```
int flags;
/* Get the file status flags and file access modes */
if ( (flags = fcntl (fd, F_GETFL, 0)) < 0)
    err_sys("F_GETFL error");
/* Set a socket as nonblocking */
if (fcntl(fd, F_SETFL, flags | O_ASYNC | O_NONBLOCK) < 0)
    err_sys("F_SETFL error");
```

- Turn off asynchronous I/O mode

```
int flags;
/* Get the file status flags and file access modes */
if ( (flags = fcntl (fd, F_GETFL, 0)) < 0)
    err_sys("F_GETFL error");
/* Turn off non-blocking mode on socket */
if (fcntl(fd, F_SETFL, flags & ~O_ASYNC & ~O_NONBLOCK) < 0)
    err_sys("F_SETFL error");
```

Signal-driven I/O: use `ioctl()`

```
int on = 1;
/* Set a socket as nonblocking */
ioctl(fd, FIOASYNC, (char *)&on);
ioctl(fd, FIONBIO, (char *)&on)
on = 0;
/* Turn off non-blocking mode on socket */
ioctl(fd, FIOASYNC, (char *)&on);
ioctl(fd, FIONBIO, (char *)&on)
```

SIGIO on sockets

- UDP socket: The signal SIGIO is generated whenever
 - A datagram arrives for the socket
 - An asynchronous error occurs on the socket
- TCP socket: the following conditions all cause SIGIO to be generated(very complex)
 - A connection request has completed on a listening socket
 - A disconnect request has been initiated
 - A disconnect request has completed
 - Half of a connection has been shut down
 - Data has arrived on a socket
 - Data has been sent from a socket (i.e., the output buffer has free space)
 - An asynchronous error occurred

Example: signal-driven I/O on UDP socket

- See source code

ADVANCED I/O FUNCTIONS

Tong Van Van, SoICT, HUST

Socket Timeouts

- There are three ways to place a timeout on an I/O operation involving a socket:
 - Call alarm, which generates the SIGALRM signal when the specified time has expired
 - Block waiting for I/O in *select*
 - Use the newer SO_RCVTIMEO and SO_SNDTIMEO socket options
- Timeout on *connect* operation?

connect with a timeout

```
#include <signal.h>
typedef void sigfunc(int)
static void connect_alarm(int);
int connect_timeo(int sockfd, const SA *saptr, socklen_t salen,
                  int nsec)
{
    sigfunc *sigfunc;
    int n;
    sigfunc = signal(SIGALRM, connect_alarm);
    if (alarm(nsec) != 0)
        err_msg("connect_timeo: alarm was already set");
    if ((n = connect(sockfd, saptr, salen)) < 0) {
        close(sockfd);
        if(errno == EINTR)
            errno = ETIMEDOUT; }
    alarm(0); // turn off the alarm
    signal(SIGALRM, sigfunc); //restore previous signal handler
    return (n);
}
static void connect_alarm(int signo) {return;}
```

readv() and writev() Functions

```
#include <sys/uio.h>
```

```
ssize_t readv(int sockfd, const struct iovec *iov, int iovcnt);  
ssize_t writev(int filedes, const struct iovec *iov, int iovcnt);
```

- Arguments:
 - *iov*: a pointer to an array of iovec structures
 - *iovcnt*: number of elements in iov array
- iov structure

```
struct iovec {  
    void *iov_base;  
    size_t iov_len;  
};
```

Example

```
#include <sys/types.h>
#include <sys/uio.h>
#include <unistd.h>
//...
ssize_t bytes_read;
int fd;
char buf0[20];
char buf1[30];
char buf2[40];
int iovcnt;
struct iovec iov[3];
iov[0].iov_base = buf0;
iov[0].iov_len = sizeof(buf0);
iov[1].iov_base = buf1;
iov[1].iov_len = sizeof(buf1);
iov[2].iov_base = buf2;
iov[2].iov_len = sizeof(buf2);
//...
bytes_read = readv(fd, iov, 3);
//...
```

recvmsg () and sendmsg ()

```
#include <sys/socket.h>
```

```
ssize_t recvmsg(int sockfd, struct msghdr *msg, int flags);
```

```
ssize_t sendmsg(int sockfd, struct msghdr *msg, int flags);
```

- Arguments:

- *msg*: pointer to msghdr structures

```
struct msghdr {  
    void *msg_name; /* protocol address */  
    socklen_t msg_namelen; /* size of protocol address */  
    struct iovec *msg_iov; /* scatter/gather array */  
    int msg_iovlen; /* # elements in msg_iov */  
    void *msg_control; /* ancillary data (cmsghdr struct) */  
    socklen_t msg_controllen; /* length of ancillary data */  
    int msg_flags; /* flags returned by recvmsg() */  
};
```

Example

```
#include <sys/socket.h>
struct sockaddr_in dest;
int rc;
struct iovec iov[3];
struct msghdr mh;
memset(&dest, '\0', sizeof(dest)); dest.sin_family = AF_INET;
memcpy(&dest.sin_addr, host->h_addr, sizeof(dest.sin_addr));
dest.sin_port = htons(TRANSACTION_SERVER);
iov[0].iov_base = (caddr_t)head; iov[0].iov_len = sizeof(struct header);
iov[1].iov_base = (caddr_t)trans; iov[1].iov_len = sizeof(struct record);
iov[2].iov_base = (caddr_t)trail; iov[2].iov_len = sizeof(struct trailer);
mh.msg_name = (caddr_t) &dest; mh.msg_namelen = sizeof(dest);
mh.msg_iov = iov; mh.msg_iovlen = 3;
mh.msg_control = NULL;
mh.msg_controllen = 0;
rc = sendmsg(s, &mh, 0); /* no flags used */
```